

Lightweight Policy-Space Control for Physics-Based Characters

ANONYMOUS AUTHOR(S)
SUBMISSION ID: PAPERS_1462

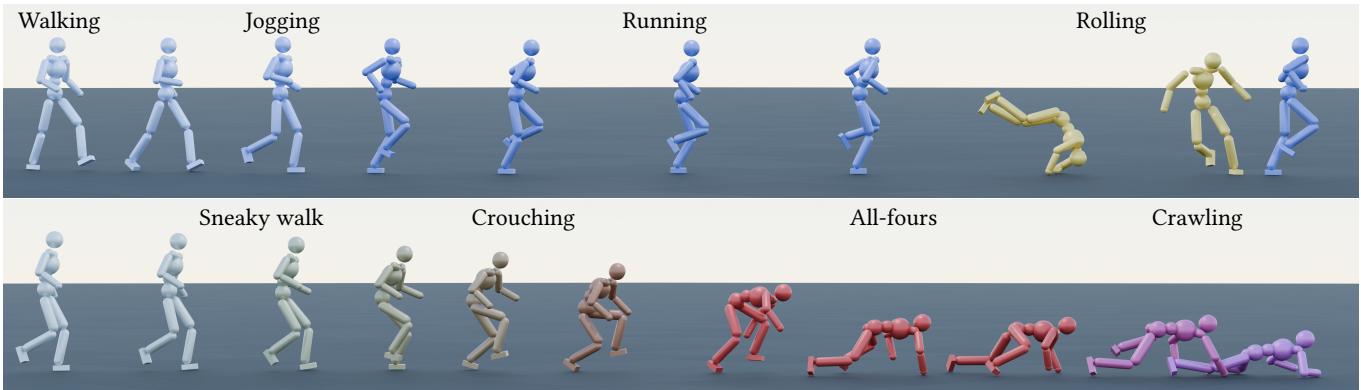


Fig. 1. Our lightweight control strategy can offer a full move-set of physics-based simulations with a continuum of real-time variations. Policies with similar state spaces can be linearly interpolated, while the systems still allow for animating episodic motions as Rolling. These motions are controlled by a gamepad with minimal logic

Physics-based character animation often relies on large policy networks or complex frameworks, which demand significant computational resources. These systems typically require partial or complete retraining to integrate new motions, limiting their flexibility during prototyping. We propose a novel approach that uses a lightweight network and dynamically changes weights, enabling new behaviors with minimal complexity. Our technique animates characters through simple operations such as policy-space linear interpolation, policy switch, and graph search. Central to our method is a weight-space regularization strategy, which organizes pre-trained policies to allow seamless transitions between motions and diverse animation styles. This approach supports real-time control of character animations while maintaining low computational overhead, preserving the benefits of physics-based controllers. Moreover, new policies can be incorporated into the system without retraining existing components. Results demonstrate that interpolated policies produce smoother transitions than direct policy switch while maintaining high motion quality comparable to state-of-the-art imitation-based frameworks at a significantly lower computational cost. An interesting finding is an equivalence between our interpolated policies and direct action interpolation, opening research paths to clarify how the policies encode actions. The possibility of interpolating single layers resulting in motion variations of the internal pattern of the motion, suggests that ours is a powerful representation that can open the path to improved control over animation variability. We showcase an application where gamepad input drives desired motions with minimal logic, highlighting the system's real-time responsiveness and simplicity.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM 0730-0301/2025/1-ART
<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

CCS Concepts: • Computing methodologies → Physical simulation; Motion processing.

Additional Key Words and Phrases: Control Policies, Deep Reinforcement Learning, Interactive Animation, Weight-space

ACM Reference Format:

Anonymous Author(s). 2025. Lightweight Policy-Space Control for Physics-Based Characters. *ACM Trans. Graph.*, 1, 1 (January 2025), 6 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 INTRODUCTION

Physics-based character animation has become a cornerstone of modern computer graphics, enabling the creation of highly realistic and physically plausible motion. This type of animation allows for interaction with the environment calculating forces and contacts in a physics simulator. A deep reinforcement learning policy is usually capable of learning complex motor skill. However, learning new skills often leads to catastrophic forgetting of the previous ones. Researchers have been working on methods to learn articulated sequences of skills, as the simple interpolation of action produced by different policies does not always keep the realism of the trained motion. Many methods resort to a two-level structure where a high-level model composes the effects of the low-level policies to achieve more complex skills. This usually entails having to retrain, partially or totally, the model when a new motion is added to the set. These methods often rely on training latent encoding of the motion set or other generative modeling setups that require heavy computation resources.

We propose a novel control technique for efficiently animating characters using compact policy networks. Our approach emphasizes simplicity and computational efficiency while preserving the robustness and adaptability of physics-based systems. By leveraging a lightweight network our method generates diverse motion

115 styles with minimal computational overhead. A key innovation is
 116 our weight-space regularization strategy, which organizes a set of
 117 pre-trained, policies into a structured framework. This allows seam-
 118 less transitions between motion types, reduces computational costs,
 119 and enables the addition of new policies without retraining existing
 120 components.

121 Our method retains the physical plausibility and robustness of
 122 physics-based animation, ensures smoother motion transitions com-
 123 pared to naive policy switching, and remains compatible with resource-
 124 limited systems, making it ideal for applications like video games
 125 and virtual reality. Experimental results demonstrate its effective-
 126 ness in animating characters with minimal logic, such as gamepad
 127 control, while achieving visually equivalent motions to those pro-
 128 duced by action interpolation. Our representation, however, allows
 129 us to manipulate the motion generation by changing the weights,
 130 introducing variability to the internal motion patterns modifying
 131 how the actions are generated.

133 2 RELATED WORK

134 Our work has been inspired by the recent development in the field of
 135 physics-based animation. We identify 3 broad areas of overlap: The
 136 composition of different motor skills from policy-based controllers,
 137 smooth transition among different motions, and directable control
 138 of the animations.

140 2.1 Composing Motor Skills

141 Motion imitation policies for character animation have traditionally
 142 been limited to reproducing a single complex motor skill. Even
 143 with large neural networks, catastrophic forgetting has hindered
 144 the ability to learn a large-scale motion dataset with a single policy.
 145 The approaches proposed by [Peng et al. 2018] introduce a one-hot
 146 conditional policy for selecting which skill to imitate, or employ
 147 the value function to sample an appropriate policy from the current
 148 state among a set of pre-trained policies. Goal inputs allow to learn
 149 high-level task.

150 Ho and Ermon [2016] employs Generative Adversarial Imitation
 151 Learning (GAIL) to achieve high-quality, realistic motor skills by
 152 training a generator to produce actions that mimic expert demon-
 153 strations, refined through an adversarial training process.

154 In the same vein, Peng et al. [2021] explores Adversarial Motion
 155 Priors (AMP) for learning motion skills from unstructured datasets
 156 by dynamically selecting and interpolating motions without explicit
 157 clip selection or sequencing. AMP effectively mitigates catastrophic
 158 forgetting and enables the composition of various motor skills from
 159 diverse datasets. Here a discriminator loss is combined with a goal
 160 loss for leveraging imitation and high-level tasks.

161 Peng et al. [2022] developed ASE (Adversarial Skills Embedding)
 162 to separate the motion imitation task from the goal task. First, it
 163 learns a low-level imitation policy on the full dataset alongside a
 164 latent state representation. Novel goal-based tasks can be trained
 165 using a high-level controller that reuses the action set learned by the
 166 low-level policy through the latent encoding. This method confines
 167 the set of poses used for solving the task to a set of highly realistic
 168 poses, with the drawback of the pose set being discrete.

172 2.2 Smooth transitions in a Policy-Based Setting

173 Motion blending is a cornerstone of creating smooth and realistic
 174 character animations, whether in classical kinematic settings or
 175 modern physics-based frameworks. In kinematic-based animation,
 176 techniques such as barycentric interpolation, radial basis functions,
 177 k-nearest neighbors, and inverse blending optimization have been
 178 extensively analyzed by Feng et al. [2012]. Their study provides a
 179 detailed comparison of these methods, highlighting their respective
 180 strengths and limitations in achieving fluid transitions.

181 Modern advancements have leveraged machine learning and
 182 physics-based approaches to push the boundaries of motion blend-
 183 ing. For instance, Juravsky et al. [2022] introduced a system that
 184 integrates natural language processing (NLP) with physics-based
 185 character control. It enables users to command character motions
 186 using natural language, bridging the gap between intuitive user
 187 input and complex motion generation.

188 Tessler et al. [2024] presented a unified system for physics-based
 189 control that employs masked motion in-painting. This approach gen-
 190 erates diverse motions from partial descriptions, such as keyframes
 191 or text commands, offering greater flexibility and creative possibili-
 192 ties in animation design.

193 Further progress was demonstrated by Luo et al. [2024a], who
 194 developed a foundational model for physics-based control. Their
 195 framework learns a universal motion representation from large, un-
 196 structured datasets and employs hierarchical reinforcement learning
 197 to produce stable, diverse human motions adaptable to various tasks.

198 Despite these advancements, challenges persist in achieving seam-
 199 less integration and realism, particularly in dynamic and interactive
 200 environments.

202 2.3 Directable Control of Physics-Based Animation

203 Achieving directable control in physics-based animation is a com-
 204 plex challenge that has seen remarkable advancements. Bergamin
 205 et al. [2019] developed a method that employs motion-matching
 206 [Clavet 2016] examples to train deep reinforcement learning policies
 207 for controllable locomotion of physically simulated characters in
 208 real-time, demonstrating its feasibility in dynamic environments.
 209 Wang et al. [2020] developed universal neural controller capable of
 210 mastering thousands of motions with different styles, supporting
 211 real-time interactive applications. Dou et al. [2023] expands the for
 212 from Peng et al. [2022] to control real-time animation by dividing
 213 skill motions into homogeneous subsets for training, offering direct
 214 control over character skills through low-level conditional models.
 215 This enables seamless and responsive character control, benefitting
 216 interactive animations. The research by [Tessler et al. 2023] enables
 217 user-controlled virtual characters by leveraging adversarial training
 218 and imitation learning. This model captures the complexity and di-
 219 versity of human motion, allowing user-controlled characters to be
 220 both directable and intuitive. The approach jointly learns a control
 221 policy and a motion encoder, enabling real-time behavior control
 222 and motion conditioning for high-level task performance. The work
 223 by Luo et al. [2024b] presents a comprehensive motion represen-
 224 tation that covers a wide range of motor skills, using an encoder-
 225 decoder structure to distill skills from a large dataset. This approach
 226 achieves high coverage and expressiveness, making it effective for

both generative tasks and motion tracking. Another significant advancement is the work by Ren et al. [2023] uses diffusion-based human motion models to generate instruction-driven animations of physics-based characters. This framework captures complex relationships between high-level human instructions and character motions, achieving state-of-the-art results in various tasks.

2.4 This work

Our approach differs from other works by viewing the policy as a dynamic object that can change its weights to change its behavior. We train our policies using a modified implementation of Peng et al. [2021] and regularize the weights to preserve similarity for leveraging simple operations in the weight-space. Linear piecewise interpolation allows us to navigate the regularization graph to obtain smooth behavior changes. The main contributions of this work are

- A regularization strategy to preserve similarity of the weights during policy training
- A novel policy network design allowing dynamic weight manipulation to adapt and control character behavior.
- The use of simple operations, such as interpolation, policy switch, and graph search, to control character behavior with minimal gamepad logic.

3 METHODOLOGY

In this section we break down the methodology of our approach, by describing our weight regularization, training strategy, and gamepad control logic.

3.1 Policy Regularization Strategy

Our approach incorporates heuristics that enable linear interpolation within the weight space of the actor network. To achieve this, we adopt a regularization strategy that preserves policy similarity for similar motions. Specifically, we utilize a policy gradient algorithm, Proximal Policy Optimization (PPO), in combination with a Multi-Layer Perceptron (MLP) with ReLU activations. This architecture supports meaningful policies under linear interpolation, unlike other architectures, such as Gated Recurrent Units (GRUs) [Xu and Karamouzas 2021; Xu et al. 2023].

We build on the Adversarial Motion Prior (AMP) framework [Peng et al. 2021], which jointly trains a discriminator and actor network for motion imitation. To apply the same motion prior across different tasks, we incorporate a task-specific goal by defining appropriate rewards. Initially, we train a primary policy, referred to as Main policy, starting from random weights without regularization. Subsequent policies are trained with regularization by applying a Gaussian prior to the network weights, encouraging similarity to the weights of policies trained for related tasks. The overall loss function is given by:

$$\mathcal{L} \equiv c_{\text{Disc}} \mathcal{L}_{\text{Disc}} + c_{\text{Task}} \mathcal{L}_{\text{Task}} + \lambda ||w - w_{\text{Reg}}||^2,$$

where w_{Reg} represents the regularization weights, c_{Disc} the importance given to the imitation task, c_{Task} the importance given to the high-level task, and λ the penalization for the weights moving away from the reference. Here $c_{\text{Disc}} + c_{\text{Task}} = 1$ and we set $\lambda = 0.02$.

The regularization strategy follows a hierarchical graph structure, starting from Main policy and branching into derived policies based on state space overlap, as illustrated in Figure 2.

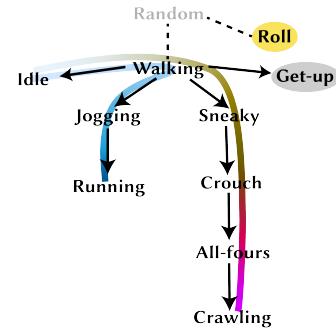


Fig. 2. Graph-based structure for the regularization strategy. Two main branches allow linear piecewise interpolation for Idle to Running, and Idle to Crawling. Random here is a different set of weights for the two deriving policies.

Policies are trained for motion imitation alongside a target localization task. In this task, a target is placed on the x-y plane, and the character is rewarded for moving toward the target and stopping upon reaching it. By maintaining a consistent input structure across policies, linear interpolation becomes feasible while preserving the character's controllability. The reward function can be specialized according to the learned skill, for example, the Idle policy is rewarded for maintaining a stationary stance, regardless of the target location.

3.2 Policy architecture and training

The policies are implemented as a multi-layer perceptron (MLP) with two fully connected hidden layers of sizes [128, 32], using ReLU activations. The input is a 225-dimensional state observation that includes the root-bone height, local body positions, rotations, velocities, angular velocities, and the relative x-y coordinates of the target location with respect to the character's root bone. The output is a 28-dimensional action vector for the Humanoid character, as described in [Peng et al. 2021]. This output represents the mean of a multivariate Gaussian distribution, with the standard deviation fixed across all dimensions. The discriminator network, also modeled as an MLP with hidden layers of sizes [1024, 512], is used during training to evaluate the imitation quality and during replay to automate the termination of episodic policies, such as Get-up or Roll motions.

To effectively learn motions it is found to be beneficial dividing the training into phases with different combinations of the parameters c_{Disc} and c_{Task} . At the beginning of the training, it is convenient to set $c_{\text{Disc}} = 1$ and consequently $c_{\text{Task}} = 0$, to allow the policy to learn the motion of quality faithful to the dataset. In this phase, it is important to tune the gradient penalty to give the possibility for complex motions to be learned. Once the motion is learned, c_{Task} can be increased to allow the high-level task to be learned. In the final phase, the policy is trained on a wavy terrain to improve its robustness to falling.

343 3.3 Gamepad Control

344 Character control is achieved by placing the x-y target in the desired
 345 direction relative to the character, where the motion speed is deter-
 346 mined by the target's distance. Switching between policies leads to
 347 abrupt behavior changes, which can be leveraged to perform specific
 348 motion skills. For instance, rolling or getting up from the floor can
 349 be executed for a fixed number of frames upon receiving an input
 350 command. The discriminator trained for the Walk motion evaluates
 351 the character's current pose and determines when it aligns with
 352 the distribution of walking states, allowing seamless transitions
 353 back to locomotion. The threshold for the discriminator is chosen
 354 empirically to be triggered from sufficiently in-distribution states.
 355

356 Our regularization approach enables linear interpolation to smoothly
 357 transition between policies. This ensures the character remains "in-
 358 distribution" during transitions by leveraging the robustness of
 359 physics-based policies. Linear piecewise interpolation along the reg-
 360 ularization graph achieves smooth, controllable behavior, avoiding
 361 the abrupt changes caused by direct policy switch. Simple policy
 362 switch, by contrast, often results in frequent falls and unrealistic
 363 motion transitions.

364 To run our control experiments, we set up a gamepad logic to
 365 control the character's behavior using Isaac Sim. We map the con-
 366 troller input from the left stick to control the character's direction
 367 and we assign the left and right analog triggers to linear piecewise
 368 interpolation paths to control the motion style. Other buttons are
 369 programmed to perform occasional policy switch as for Roll and
 370 Get-up that are eventually terminated by a threshold on the Walk
 371 discriminator.

372 To ensure real-time performance, all policies are preloaded to
 373 eliminate file-loading delays. Policy decisions, linear interpolation,
 374 and policy evaluation occur seamlessly between frames without
 375 impacting the framerate.

376 4 RESULTS

377 The results demonstrate that our control strategy enables smooth
 378 transitions between policies, effectively bridging the gap between
 379 imitation policies.

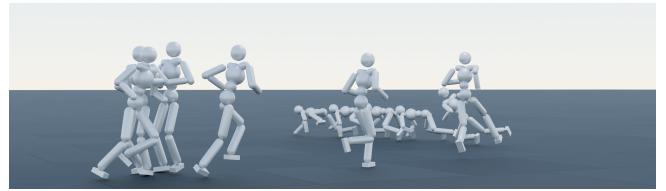
380 Figure 1 illustrates how motion is smoothly interpolated across
 381 different styles and how episodic policies are seamlessly integrated
 382 into the control framework. The linear piecewise interpolation fol-
 383 lows a path on the policy graph shown in Figure 2. During inference,
 384 the interpolation of 33K parameters and the evaluation of a single
 385 network [128,32] collectively result in approximately 100K multiply-
 386 add operations per simulation step.

388 4.1 Comparison with Action Interpolation

389 Interestingly, the visual quality of the motion generated by our
 390 method is equivalent to that produced by interpolating the actions
 391 from the same two networks. During inference, evaluating two
 392 networks [128,32] and interpolating 28 action values amount to
 393 approximately 64K multiply-add operations per simulation step.
 394

395 4.2 Comparison with ASE

396 To compare our method with ASE, we trained a low-level imitation
 397 policy using the same dataset employed for our set of regularized
 398



400 Fig. 3. With the ASE framework we obtain controllable motion of high
 401 quality, however, it can only use the actions learned from the low-level
 402 controller to perform the high-level task. In this example, the motion is
 403 controllable only when the character is standing, from the moment it falls,
 404 the character starts imitating crawling motions and becomes indifferent to
 405 the control input

406 policies, excluding motion for the "Roll" animation to prevent acci-
 407 dental triggering during high-level tasks. To ensure a fair compari-
 408 son, we aimed to provide control over the character's speed and
 409 z-level of the root bone, aligning with our selection of motions in
 410 our policy graph.

411 We defined an augmented target location high-level task with
 412 four goal inputs: the x-coordinate and y-coordinate of the target
 413 location, the root-bone speed, and the z-level of the root bone. The
 414 character was rewarded for achieving the desired location, speed,
 415 and z-level using a reward function of the form:

$$r \equiv r_w e^{-r_s(v-v_{\text{desired}})^2},$$

416 where v represents the parameter being compared to the desired
 417 value v_{desired} , r_s controls the exponential decay scale as the parame-
 418 ter deviates from the desired value, and r_w weights the parameter's
 419 contribution to the total reward.

420 The resulting motion is shown to be controllable for reaching
 421 the target location; however, the inputs for speed and z-level were
 422 largely ignored by the high-level task. As demonstrated in figure
 423 3 and the supplementary video, the character could walk and run
 424 toward the desired direction but eventually fell and started crawling,
 425 rendering the location task infeasible due to insufficient poses for
 426 crawling. The network largely disregarded the speed and z-level
 427 inputs. This experiment highlighted ASE's ability to reproduce high-
 428 quality motions using a discrete set of actions from the dataset.
 429 However, with a small dataset, our method offers greater motion
 430 variety. As a low-level imitation policy guided by high-level goals,
 431 our method generates novel actions, effectively overcoming dataset
 432 limitations.

433 During inference, ASE evaluates three networks: the high-level
 434 policy [1024,512], a style network [512,256] transforming the latent z ,
 435 and the low-level policy [1024,1024,512]. Accounting for activations,
 436 this amounts to approximately 2.2M multiply-add operations per
 437 simulation step, 22 times more than our method and 34 times more
 438 than action interpolation. We have trained the low-level policy for
 439 40K epochs and the high-level task for 32K epochs with 4096 parallel
 440 environments.

443 4.3 Interpolation of single layers

444 The results in the supplementary video show that interpolation of a
 445 single hidden layer or the output layer results in motion variations

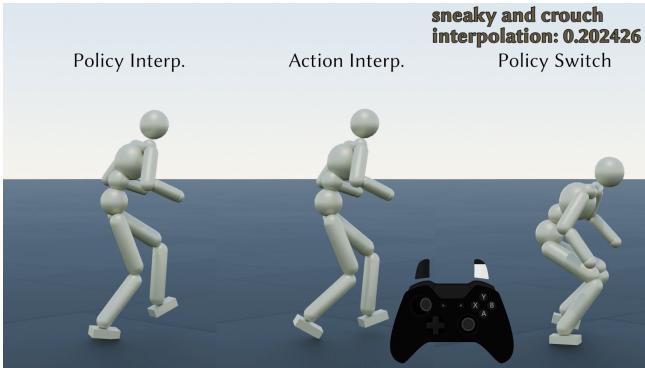


Fig. 4. An illustrative frame from the supplementary video. Our method produces poses equivalent to direct action interpolation, while direct policy switching fails to ensure smooth transitions.

that do not directly translate into an interpolation of the final action. Instead, these variations indicate that weight interpolation affects the internal representations in a way that alters the motion trajectory, producing behaviors that deviate from simple action blending. This suggests that the intermediate layers of the network encode complex, nonlinear transformations that capture the temporal and spatial dynamics of the motion.

The observed phenomenon highlights the role of the network's architecture in shaping the interpolation outcomes. Interpolating weights in hidden layers influences the latent structure of the policy, which can lead to transitions that are not easily predictable from the final action alone. These results underline the potential of weight manipulation as a tool for generating diverse and expressive motion patterns beyond what is achievable through action interpolation.

5 DISCUSSION

Our method demonstrates the effectiveness of weight regularization for preserving similarity during policy training, enabling interpolation in the weight space. As shown in Figures 5 and 6, regularized policies with similar state distributions cluster closely in PCA representation, while non-regularized policies or those with highly dissimilar state distributions are farther apart.

However, regularization can sometimes slow or limit the training of specific skills. Using the AMP algorithm, our regularization is effective when the state spaces of the motions are similar. For instance, training a chain of policies from "Walking" to "Crawling" requires approximately 2,000 epochs with 4,096 environments running in parallel for each regularized training session. However, regularizing a jump motion from idle or another locomotion policy can hinder training, as the RL algorithm focuses on matching the jump's starting and landing states. We attribute this limitation to jumps requiring significantly higher torque than locomotion. Encouraging jumps by rewarding states with higher Z-positions of the root bone could help, yet it may reduce realism, as the policy might adopt non-imitative strategies to achieve the goal.

The supplementary video demonstrates that our method achieves visual equivalence to direct action interpolation. Figure 4 highlights

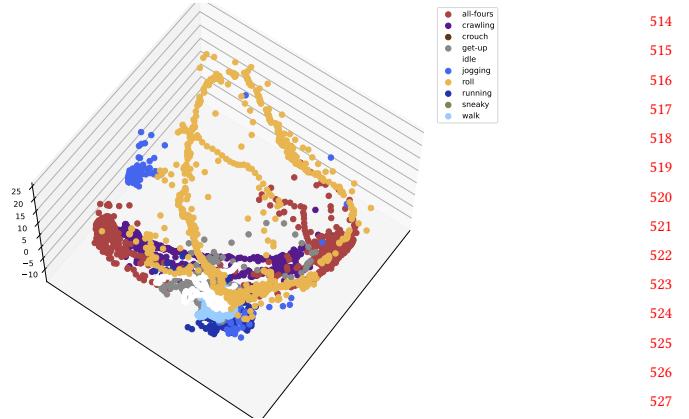


Fig. 5. 3D PCA representation of states visited by the policies. Regularized policies produce realistic transitions for motions close in state space.

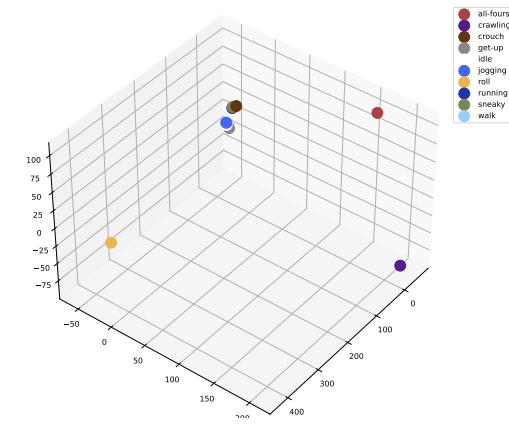


Fig. 6. 3D PCA representation of policy weights. Policies close in this reduced space yield realistic results during interpolation. Comparison with Figure 5 shows the correlation between weight distances and state distribution overlap.

how the same gamepad input produces consistent poses. This behavior may stem from the structure of the policy networks, which consist of linear layers with ReLU activations. While ReLU introduces localized nonlinearity by zeroing out negative values, the interpolation remains piecewise linear within regions where the same set of neurons is active.

This piecewise linearity suggests that interpolation operates in regions where linearity is preserved. Additionally, the learned policy manifold may align so that intermediate states during interpolation remain consistent with valid actions, influenced by PPO training and data constraints. While this highlights the robustness of the network design, further investigation is needed to confirm whether this behavior generalizes across architectures or tasks.

Figure 5 shows the state distribution in PCA-reduced space, where action interpolation creates realistic transitions when end-point policies share similar state distributions. This requirement aligns with

our regularization strategy. Although action interpolation is computationally cheaper and requires less training, our representation contains more information than its output actions.

The weights w of a control policy can be interpreted as a latent motion encoding, capturing information about the training dataset's visited states, gravity's influence, and robustness to perturbations. During evaluation, this information is "collapsed" into a single action, representing the instantaneous policy output. This behavior reflects the realization of a trajectory within a high-dimensional policy manifold. Working at the weight level modifies the internal patterns of the motion, as demonstrated by the interpolation of single layers.

The observed linear characteristics during weight interpolation suggest that the policy manifold has an underlying structure that could support dimensionality reduction. This would enable more efficient and interpretable policy representations. Further analysis is required to understand how this structure generalizes across tasks and architectures and how it can be exploited for practical applications.

6 CONCLUSION

We propose a novel approach to physics-based character animation that has the potential to establish a new paradigm for representing complex human motion in simulation. By representing policies as dynamic objects capable of altering their behavior through weight manipulation, we enable variations across the full motion representation rather than restricting them to contextual actions. Our results demonstrate that interpolating individual layers within the policies generates motion variations, showing that weight manipulation is more expressive than simply blending output motions. This reveals the ability to uncover novel motion trajectories that are inaccessible through traditional action interpolation techniques.

Thus far, this approach has been effectively applied to motions with similar state distributions, consistent with the conditions under which action interpolation achieves realistic results. However, we see promising opportunities to extend this method to a broader range of motions with more diverse state distributions. This presents a chance to explore new policy representations that enhance the effectiveness and interpretability of weight manipulation, transitioning from "black-box" systems to a more explainable and versatile representation of motion.

Future research should focus on developing techniques to generalize this approach to more diverse motion datasets while ensuring smooth transitions between motions with dissimilar dynamics. Incorporating domain knowledge could further refine weight manipulation strategies and improve guidance. Additionally, systematic analysis of how specific weights contribute to motion variations could lead to more precise control over motion style and behavior. Understanding the structure of the policy manifold could also support dimensionality reduction, optimizing storage, computation, and policy interpretation.

Another avenue for exploration is incorporating learning objectives that encourage interpretable latent encodings, fostering a deeper understanding of the relationship between weights, dynamics, and the resulting motion. Finally, integrating these techniques into interactive or real-time systems could unlock applications in

animation, robotics, and virtual environments, where adaptability and explainability are critical.

REFERENCES

- Kevin Bergamin, Simon Clavet, Daniel Holden, and James Richard Forbes. 2019. DReCon: data-driven responsive control of physics-based characters. *ACM Trans. Graph.* 38, 6, Article 206 (Nov. 2019), 11 pages. <https://doi.org/10.1145/3355089.3356536>
- Simon Clavet. 2016. Motion Matching and The Road to Next-Gen Animation. Identifier: GDC2016Clavet, Addendum: 2016-10-17, Scanner: Internet Archive HTML5 Uploader 1.6.3.
- Zhiyang Dou, Xuelin Chen, Qingnan Fan, Taku Komura, and Wenping Wang. 2023. C-ASE: Learning Conditional Adversarial Skill Embeddings for Physics-based Characters. In *SIGGRAPH Asia 2023 Conference Papers* (Sydney, NSW, Australia) (*SA '23*). Association for Computing Machinery, New York, NY, USA, Article 2, 11 pages. <https://doi.org/10.1145/3610548.3618205>
- Andrew Feng, Yazhou Huang, Marcelo Kallmann, and Ari Shapiro. 2012. An Analysis of Motion Blending Techniques. In *Motion in Games*. 232–243.
- Jonathan Ho and Stefano Ermon. 2016. Generative Adversarial Imitation Learning. *CoRR* abs/1606.03476 (2016), arXiv:1606.03476 <http://arxiv.org/abs/1606.03476>
- Jordan Juravsky, Yunrong Guo, Sanja Fidler, and Xue Bin Peng. 2022. PADL: Language-Directed Physics-Based Character Control. In *SIGGRAPH Asia 2022 Conference Papers* (*SA '22*). ACM, 1–9. <https://doi.org/10.1145/3550469.3555391>
- Zhengyi Luo, Jinkun Cao, Josh Merel, Alexander Winkler, Jing Huang, Kris M. Kitani, and Weipeng Xu. 2024a. Universal Humanoid Motion Representations for Physics-Based Control. In *The Twelfth International Conference on Learning Representations*. <https://openreview.net/forum?id=OrOd8PxOO2>
- Zhengyi Luo, Jinkun Cao, Josh Merel, Alexander Winkler, Jing Huang, Kris M. Kitani, and Weipeng Xu. 2024b. Universal Humanoid Motion Representations for Physics-Based Control. In *The Twelfth International Conference on Learning Representations*. <https://openreview.net/forum?id=OrOd8PxOO2>
- Xue Bin Peng, Pieter Abbeel, Sergey Levine, and Michiel van de Panne. 2018. DeepMimic: example-guided deep reinforcement learning of physics-based character skills. *ACM Trans. Graph.* 37, 4, Article 143 (July 2018), 14 pages. <https://doi.org/10.1145/3197517.3201311>
- Xue Bin Peng, Yunrong Guo, Lina Halper, Sergey Levine, and Sanja Fidler. 2022. ASE: large-scale reusable adversarial skill embeddings for physically simulated characters. *ACM Trans. Graph.* 41, 4, Article 94 (July 2022), 17 pages. <https://doi.org/10.1145/3528223.3530110>
- Xue Bin Peng, Ze Ma, Pieter Abbeel, Sergey Levine, and Angjoo Kanazawa. 2021. AMP: adversarial motion priors for stylized physics-based character control. *ACM Trans. Graph.* 40, 4, Article 144 (July 2021), 20 pages. <https://doi.org/10.1145/3450626.3459670>
- Jiawei Ren, Mingyuan Zhang, Cunjun Yu, Xiao Ma, Liang Pan, and Ziwei Liu. 2023. InsActor: instruction-driven physics-based characters. In *Proceedings of the 37th International Conference on Neural Information Processing Systems* (New Orleans, LA, USA) (*NIPS '23*). Curran Associates Inc., Red Hook, NY, USA, Article 2618, 13 pages.
- Chen Tessler, Yunrong Guo, Ofir Nabati, Gal Chechik, and Xue Bin Peng. 2024. MaskedMimic: Unified Physics-Based Character Control Through Masked Motion Inpainting. *ACM Trans. Graph.* 43, 6, Article 209 (Nov. 2024), 21 pages. <https://doi.org/10.1145/3687951>
- Chen Tessler, Yoni Kasten, Yunrong Guo, Shie Mannor, Gal Chechik, and Xue Bin Peng. 2023. CALM: Conditional Adversarial Latent Modelsnbsp; for Directable Virtual Characters. In *ACM SIGGRAPH 2023 Conference Proceedings* (Los Angeles, CA, USA) (*SIGGRAPH '23*). Association for Computing Machinery, New York, NY, USA, Article 37, 9 pages. <https://doi.org/10.1145/3588432.3591541>
- Tingwu Wang, Yunrong Guo, Maria Shugrina, and Sanja Fidler. 2020. UniCon: Universal Neural Controller For Physics-based Character Motion. arXiv:2011.15119 [cs.GR]
- Pei Xu and Ioannis Karamouzas. 2021. A GAN-Like Approach for Physics-Based Imitation Learning and Interactive Character Control. *Proc. ACM Comput. Graph. Interact. Tech.* 4, 3, Article 44 (Sept. 2021), 22 pages. <https://doi.org/10.1145/3480148>
- Pei Xu, Xiumin Shang, Victor Zordan, and Ioannis Karamouzas. 2023. Composite Motion Learning with Task Control. *ACM Trans. Graph.* 42, 4, Article 93 (July 2023), 16 pages. <https://doi.org/10.1145/3592447>